

HEMANTH K R

+91 8939565145 • ktrhemanthh@gmail.com • Chennai, India • [Linkedin](#) • [Github](#) • [Website](#)

SKILLS

Data Engineering: Apache Airflow, Airbyte, Mage AI, PostgreSQL, AWS Athena, AWS Glue

Cloud & Infrastructure: AWS, GCP, Docker, Grafana, GCP Cloud Monitoring

Machine Learning & AI: PyTorch, TensorFlow, LangGraph, Agentic AI, Computer Vision, RAG, YOLO, Fine-tuning, Transformers, scikit-learn, Pandas

Frameworks: FastAPI, Node.js, Django, React

Languages: Python, SQL, Go, Java, C++

EDUCATION

B.Tech in Artificial Intelligence and Data Science

Sep 2022 – May 2026

Sri Venkateswara College of Engineering, **CGPA: 8.1/10**

WORK EXPERIENCE

Third I

AI Data Engineer

Jul 2026 - Present

- Played a key role in migrating AI and data platform infrastructure from **AWS to GCP**, re-platforming compute, storage, and database workloads while ensuring **zero downtime and zero data loss**
- Designed and developed a **data observability platform** using Grafana, Prometheus, and GCP Cloud Monitoring to monitor **15–20 production pipelines**, tracking health, errors, logs, and data consistency through Grafana Infinity dashboards
- Managed GCP data platform infrastructure, troubleshooting production issues and improving the reliability of data workloads

AI Data Engineer Intern

Jul 2025 – May 2026

- Engineered **15 ETL pipelines** using Airbyte, Mage AI, Python, PostgreSQL and AWS S3, processing over **100,000 advertising records** daily from multiple marketing APIs
- Designed **medallion-based** data models for marketing analytics, streamlining reporting and downstream analysis workflows
- Enhanced analytical query efficiency by **20%** by using **schema tuning** and implementing **indexing strategies**
- Collaborated with stakeholders to translate **reporting requirements** into **scalable data pipelines**, supporting marketing performance analysis and decision-making

Hyundai Motors

AI Intern

Jan 2023 - Feb 2023

- Established an end-to-end data pipeline for **2000 image samples**, supporting preprocessing, augmentation, and feature extraction for computer vision models
- Developed a **machine learning workflow** with a Tkinter-based interface, supporting dataset management, model selection, training orchestration, and metric visualisation
- Increased YOLO validation accuracy from **80% to 85%** by optimising model hyperparameters and applying data augmentation techniques

PROJECTS

Multimodal Image Search & Asset Management Platform

May 2026

- Developed a multimodal image search platform for a client, enabling natural language retrieval across **1000+ production creative assets** using Gemini multimodal embeddings and PostgreSQL (pgvector)
- Built a FastAPI backend to asynchronously process image uploads, metadata extraction, and embedding generation using FastAPI background tasks
- Developed a React-based asset management interface featuring semantic search, hierarchical folder organisation, and image editing workflows
- Integrated the **SAM (Segment Anything Model)** to enable interactive object segmentation and background removal for selected image regions

CERTIFICATIONS

PyTorch and Deep Learning for Decision Makers (LFS116) | Linux Foundation | [Credential](#)

Databricks Fundamentals Accreditation | Databricks | [Credential](#)

Python Gold Badge and SQL Intermediate Certification | HackerRank | [Credential](#)

Google Cybersecurity Professional Certificate | Google | [Credential](#)